

Ok let's focus on these two parts:

pose\_estimation and object\_detection - classification.

Everything that's done right now and everything left and everything else needed. give me directly.

Thought for 1m 6s >

Thought for 1m 40s >

Got it—here's the straight status for the two pieces, no fluff.

## Pose estimation

### What's already in the repo

- **Model assets folder:** aixavier/models/pose/
  - `download.sh` (grabs `yolov8n-pose.pt`).
  - `export_trt.sh` (exports ONNX and builds **FP16** TensorRT: `engines/yolov8n-pose-fp16.engine`).
  - `calibrate.sh` (INT8 calibration **stub**, not complete).
  - `README.md` with target perf on Xavier.
- **Detector config:** `configs/detectors/pose_velocity.yaml` points to `models/pose/engines/yolov8n-pose-fp16.engine`, input 640x640, interval: 3.
- **Runtime skeleton:** event-bus, runners/trackers/rules services exist and expect lines in `artifacts/*`, but there's **no actual pose inference service**.

### What's missing / left to do

- **Inference service (P0):** No `detect_pose` process that:
  - Loads the **TensorRT** engine,
  - Runs keypoint inference,
  - Decodes YOLOv8-pose outputs (17 keypoints) to skeletons,
  - Publishes detections to the FileEventBus (e.g., `artifacts/detections/events.log`).

- **Velocity/analytics hook** (P0): `pose_velocity.yaml` implies velocity computation, but there's **no code** that:
  - Tracks keypoints per ID (needs tracker integration),
  - Computes per-joint velocity and emits rule-friendly events/attributes.
- **INT8 calibration** (P1): `calibrate.sh` is incomplete (... in command), needs dataset path & full `trtexec` flags.
- **Dependencies** (P1): Export scripts use `ultralalytics` and `trtexec`, but `ultralalytics` **isn't in** `requirements*.txt`; TensorRT runtime libs aren't wired in a detector image/Dockerfile.
- **Docker/service wiring** (P1): No `src/detectors/pose` code, Dockerfile, or `docker-compose service` for pose.
- **Tests** (P2): No unit/integration tests for pose decode, accuracy, or latency.

### What else you'll need to ship

- **Tracker**: ByteTrack (or KPT-aware tracker) to maintain person IDs -> stable keypoint streams.
- **Schema**: Event JSON for pose (keypoints, conf, bbox, track\_id, ts) and velocity features.
- **Perf targets**: Baseline FPS/latency on target SKUs; watchdog + metrics (Prometheus).
- **Edge packaging**: A detector container with CUDA/TensorRT, healthcheck & logs.

### Fastest path to green (pose)

1. **Create detector service** `src/detectors/pose/main.py`:
  - Load `models/pose/engines/yolov8n-pose-fp16.engine` with TensorRT.
  - Pre/postprocess to emit: `{type:"pose", camera_id, track_id, bbox, keypoints:[(x,y,conf)*17], ts}` to `artifacts/detections/events.log`.
2. **Add tracker**:
  - Start with bbox-level ByteTrack on person boxes inferred from keypoints (or run a light person detector at `interval: N` to seed IDs).
3. **Velocity module**:
  - Compute per-joint  $\Delta$  and magnitude per frame; publish `{type:"pose_velocity", joints:{...}, track_id, ts}`.
4. **Finish INT8**:

- Fix `calibrate.sh` with full `trtexec` flags and `DATASET_PATH_POSE_CALIB`, produce `engines/yolov8n-pose-int8.engine`.

## 5. Wire it up:

- Add Dockerfile, service in `docker-compose.yml` (`detect_pose`), envs, volumes.

## 6. Add tests:

- Unit test ONNX → decoder on a small clip; golden JSON for 10 frames.

# Object detection + classification

## What's already in the repo

- **YOLO detector assets:** `aixavier/models/yolo/`
  - `download.sh` (Ultralytics weights), `export_trt.sh` (builds **FP16**), `calibrate.sh` (INT8 done properly), `README.md`.
- **Detector config:** `configs/detectors/object.yaml` → `models/yolo/yolov8n-fp16.engine`, classes include `person`, `bag`, `suitcase`, `backpack`, `fire`, `smoke`.
- **Use-case**  
**recipe:** `configs/usecases/object_identification_classification.yaml` (PPE logic: `hard_hat`, `safety_vest`, `tool`, `missing_ppe` rule).
- **Runtime skeleton:** event bus, rules, tracker scaffolding (no real ByteTrack), exporter/recorder services.

## What's missing / left to do

- **Inference service (P0):** No `detect_object` process that runs YOLO TensorRT, decodes boxes/classes, emits events.
- **Class coverage mismatch (P0):** Use-case expects **PPE classes** (`hard_hat`, `safety_vest`, `tool`) but default `yolov8n` doesn't contain those → need **fine-tuned weights** or a second-stage classifier.
- **Tracking (P0/P1):** `tracker: bytetrack` is referenced but not implemented; `src/trackers` is a stub.
- **Classification stage (P1):** If you keep general YOLO, you'll need ROI **classifier** for PPE/tools from crops:

- No classifier model, code, or config present.
- **Dependencies & Docker (P1):** Same as pose: `ultralytics` not in requirements; no detector Dockerfile/service.
- **INT8 dataset pathing (P1):** `DATASET_PATH_OBJECT_CALIB` placeholder; needs real path and sample images.
- **Tests (P2):** No decode/accuracy tests; no end-to-end rule trigger tests tied to detections.

### What else you'll need to ship

- **Either:**
  - A **fine-tuned YOLO** (PPE classes) → simplest single-stage pipeline, or
  - **Two-stage:** General YOLO (person/tool-ish) → crop → lightweight classifier (e.g., MobileNetV3/ViT-Tiny) to label `hard_hat/safety_vest/tool/none` .
- **Label schema + thresholds:** Confidence, dwell seconds (already in use-case), NMS/IoU sane defaults.
- **ByteTrack** integration: associate boxes across frames for dwell/missing-PPE logic.
- **Evaluation harness:** mAP@.5:.95 for detector; top-1 for classifier; rule precision/recall on a small validation clip set.

### Fastest path to green (object + classification)

#### 1. Create detector service `src/detectors/object/main.py` :

- Load `models/yolo/yolov8n-fp16.engine` , preprocess 640×640, decode, publish:

```
{type:"object", camera_id, ts, detections:[{bbox, cls, conf, id(opt
```

#### 2. Add ByteTrack:

- Port a minimal ByteTrack impl; assign `track_id` to detections; publish `type:"object_tracked"` .

#### 3. Unblock PPE logic (pick one):

- **Option A (preferred for time):** drop-in **fine-tuned YOLO PPE engine**; update `configs/detectors/object.yaml` classes to `hard_hat/safety_vest/tool/person` .
- **Option B:** add `src/classifier/ppe/` :
  - Crop person head/torso/tool ROIs,

- Run classifier ONNX/TRT,
- Enrich detection with `attributes:{ppe:"hard_hat"/"none", ...}`.

#### 4. Rules wiring:

- Ensure emitted event fields  
`match configs/usecases/object_identification_classification.yaml` (e.g., `classes_required`, `actor_class:"person"`, `dwelt_seconds`).

#### 5. Finish INT8 (optional but good):

- Point `DATASET_PATH_OBJECT_CALIB` to a few hundred varied frames and run `calibrate.sh` → `engines/yolov8n-int8.engine`.

#### 6. Shipables:

- Detector Dockerfile + `docker-compose` service `detect_object`.
- Add `ultralytics` to `requirements-dev.txt` (export only) and document TensorRT/JetPack for runtime.
- Add unit tests for post-proc decode + a 30-sec demo E2E that triggers `missing_ppe`.

### Quick P0 checklist to start today

- Implement `src/detectors/object/main.py` + Dockerfile + compose service.
- Implement `src/detectors/pose/main.py` + Dockerfile + compose service.
- Add minimal ByteTrack and return `track_id`.
- Decide PPE path: **fine-tuned YOLO PPE** (swap engine) **or** add ROI classifier; wire to rules.
- Remove ... placeholders in `models/pose/calibrate.sh`; verify both calibrations run.
- Add `ultralytics` to dev reqs for export; document `trtexec` /TensorRT runtime.
- Add 2 tiny integration tests: one sample video → expected JSON events for both services.

If you want, I can stub both detector services (TensorRT runtime + post-proc) and a minimal ByteTrack in one go so you can drop in engines and run.

These are the models I am thinking of using. can you browse and tell me if there exists any better models. latest and better performance.

pose\_estimation → RTMPose-m.  
object\_detection and classification → YOLO11n or RT-DETRv2-R50.

Thought for 58s >

Here's the straight take, after checking the latest benchmarks and repos:

## Pose estimation

**Your pick:** RTMPose-m

**Call:** Still a strong default. It's ~75.8 AP on COCO keypoints with excellent throughput; among real-time top-down methods it's hard to beat for accuracy/speed on 17-kpt human pose. [arXiv](#) +1

If you want something *newer* that can scale better with many people per frame:

- **RTMO (one-stage, CVPR'24):** ~74.8 AP (RTMO-L) and very high FPS on V100; one-stage design avoids the top-down per-person cost, so it shines in crowded scenes. Accuracy is a hair under RTMPose-m but throughput in multi-person settings can be superior. [CVF Open Access](#)
- **Whole-body (133 kpts) focus: DWPose** distillation lifts RTMPose-L on COCO-WholeBody to ~66.5 AP (beats RTMPose-X teacher). Use if you truly need hands/face/feet. [arXiv](#) +1
- **YOLO-pose (YOLO11-pose):** convenient, export-friendly, but there isn't clear evidence it beats RTMPose-m on COCO-keypoints AP; pick it mainly for simpler tooling. [Ultralytics Docs](#)

**Bottom line (pose):**

- Keep **RTMPose-m** for best AP with typical person counts.

- Switch to **RTMO** if you're bottlenecked by many simultaneous people and want one-stage scaling. [CVF Open Access](#)

## Object detection (and your "classification" stage)

Your picks: YOLO11n or RT-DETRv2-R50

What's better/newer right now:

- **YOLO12n (2025)**: successor to YOLO11n; reported ~40.6 mAP (vs ~39.5 for YOLO11n) at similar tiny-model latency. If you want the *same nano footprint* but a bit more accuracy, YOLO12n is the cleanest upgrade. [Replicate](#) +2
- **RT-DETRv2-R50**: strong real-time transformer baseline around **53.1 AP @ ~108 FPS (T4)**; ~55.3 AP with Objects365 pretraining. Great if you can afford a midsize model and want a big accuracy jump over nano YOLO. [GitHub](#) +1
- **RF-DETR (2025)**: new open-source real-time DETR variant; **first real-time model to exceed 60 AP on COCO** (base sizes). If you prioritize accuracy while staying real-time on decent GPUs, this is currently the ceiling among real-time models. [GitHub](#) +1

**Classification add-on:**

If your pipeline needs attributes (e.g., PPE) beyond detector classes, keep a lightweight ROI classifier head (any small CNN/ViT). If you stay in the Ultralytics stack, the **YOLO12 classification** head is available; otherwise use a tiny EfficientNet/MobileNet. (YOLO12 family + export paths are documented and actively maintained.) [Ultralytics Docs](#)

**Bottom line (detection):**

- Need **nano** speed/size? Pick **YOLO12n** over YOLO11n. Small, simple upgrade. [Replicate](#) +1
- Need **higher AP but still real-time**? Use **RT-DETRv2-R50**; if you can go bigger and want the best real-time AP reported, evaluate **RF-DETR**. [GitHub](#) +2

**TL;DR picks**

- **Pose**: Stay **RTMPose-m** unless you're crowd-heavy → then **RTMO**. [arXiv](#) +1

- **Detection:** Upgrade YOLO11n → YOLO12n (same tiny tier, better mAP). For accuracy at real-time, go RT-DETRv2-R50, or RF-DETR if you can budget more compute. [Replicate](#) +3

If you want, I can map these to your target hardware (Jetson/CPU/T4/Orin) and suggest exact variants + TensorRT export flags and expected FPS.

Will this object detection model also do the classification??

Thought for 17s >

Short answer: **yes, but only for the classes it was trained on.**

- YOLO/RT-DETR are *detectors with a classifier head per box*. Every detection comes with a class label and score (e.g., person , car , dog ).
- If by "classification" you mean **PPE attributes** like hard\_hat , safety\_vest , tool , the **default COCO weights won't do it**—those labels aren't in COCO. You'll get person and generic objects, but **not** "wearing a hard hat vs not."

## What to do for PPE "classification"

You have two clean paths:

### A) Single-stage (recommended if you can get data/weights)

Fine-tune YOLO / RT-DETR on a PPE dataset so the detector's class set is:

person, hard\_hat, safety\_vest, tool (and whatever else you need).

**Pros:** simple runtime, fastest E2E, one engine.

**Cons:** needs labeled boxes (or ready weights).

**How you'd wire it:** the detector already "classifies" each box → your events directly contain PPE classes, no second pass.

### B) Two-stage (no custom detector training required)

1. Run a general detector (YOLO/RT-DETR) to get person boxes (and optionally head/torso crops).



2. Run a **light ROI classifier** (MobileNet/EfficientNet/ViT-Tiny or a YOLO "cls" head) on those crops to predict:

- ppe.hard\_hat: yes/no , ppe.safety\_vest: yes/no , tool: yes/no (multi-label).

**Pros:** flexible, easy to add/change attributes.

**Cons:** extra compute and code path.

## Quick recommendations for your picks

- If you stay tiny: use YOLO(11→)12n fine-tuned with PPE labels (A), or pair YOLO12n with a tiny ROI classifier (B).
- If you want higher AP and can afford it: RT-DETRv2-R50 fine-tuned (A). If you can't fine-tune now, use B with RT-DETRv2 for detection.

## Event schema tip (works for both paths)

For single-stage PPE detector:

```
{ "type":"object", "track_id":42, "cls":"hard_hat", "conf":0.91,
  "bbox": [x,y,w,h], "ts":... }
```

For two-stage:

```
{ "type":"object_tracked", "track_id":42, "cls":"person", "conf":0.89,
  "bbox": [x,y,w,h], "ts":..., "attributes": { "ppe": { "hard_hat": true,
  "safety_vest": false }, "tool": true } }
```

## Bottom line

- Yes, the detector "does classification" for its class list.
- For PPE or other attributes, you must (A) train/use a detector with those classes or (B) add an ROI classifier.

Tell me which way you want to go and your target device (Jetson/T4/CPU), and I'll pin exact model variants, export flags, and expected FPS.

Alright full implementation pipeline. How do i implement how do everything

## Thought for 22s >

Awesome—here's a complete, end-to-end implementation plan you can follow today. I'm giving you a **battle-tested default** plus variants so you can swap pieces without re-architecting.

## 0) TL;DR architecture

- **Cameras** → ingest (RTSP → frames)
- **Detectors:**
  - detect\_object (YOLO/RT-DETR TensorRT) → boxes + class + conf
  - detect\_pose (RTMPose or YOLO-pose TensorRT) → keypoints + conf
- **Tracker:** bytetrack → adds track\_id to object streams
- **(Optional) ROI classifier:** crops from tracked boxes → PPE attributes (multi-label)
- **Analytics:** dwell time / missing PPE / velocity on keypoints
- **Event Bus:** newline-delimited JSON (NDJSON) files or Kafka/Redis (start with files)
- **Exporter/Recorder/Alerts:** consume events → UI, DB, webhooks

## 1) Repos & layout

```
aixavier/  
  src/  
    detectors/  
      object/  
        main.py           # TRT runtime + postproc + publish  
        postproc.py       # decode, NMS  
        docker/Dockerfile  
      pose/  
        main.py  
        postproc.py  
        docker/Dockerfile  
    trackers/  
      bytetrack.py  
    classifier/  
      ppe/
```

```
infer.py                # multi-label ROI classifier (optional)
export.md
analytics/
  ppe_rules.py
  pose_velocity.py
bus/
  file_event_bus.py
configs/
  detectors/object.yaml
  detectors/pose.yaml
  classifier/ppe.yaml
  analytics/ppe_rules.yaml
  analytics/pose_velocity.yaml
models/
  object/
    engines/...
    onnx/...
  pose/
    engines/...
    onnx/...
  classifier/
    ppe_cls.onnx / .engine
docker-compose.yml
tests/
  data/ (images, short clips)
  e2e_object.yaml
  e2e_pose.yaml
```

## 2) Object detection (+ classification)

### 2.1 Choose your path

- **Path A (single-stage, recommended if you can fine-tune):** Train detector with your domain labels  
e.g., person, hard\_hat, safety\_vest, tool (+ anything else).  
→ One engine, simplest runtime.
- **Path B (two-stage, fastest to ship without training):**  
General detector ( person , etc.) + lightweight ROI classifier for PPE attributes

( hard\_hat yes/no , safety\_vest yes/no , tool yes/no ).

→ Add a small ONNX to classify crops.

I'll detail **both**; your wiring stays almost identical.

## 2.2 Model export

### Option 1: YOLO (nano/tiny friendly)

- **Train / load** (Ultralytics): yolo detect train or yolo detect predict as needed.
- **Export ONNX:**  
yolo export model=best.pt format=onnx imgsz=640 opset=12 dynamic=False
- **Build TensorRT (FP16):**

```
trtexec --onnx=models/object/onnx/yolo.onnx \
        --saveEngine=models/object/engines/yolo-fp16.engine \
        --fp16 --workspace=4096 --shapes=input:1x3x640x640
```

- **(Optional INT8):** prepare 500–1000 calibration images, then:

```
trtexec --onnx=.../yolo.onnx --int8 --fp16 \
        --saveEngine=.../yolo-int8.engine \
        --calib=max --calibCache=.../yolo.calib
```

### Option 2: RT-DETRv2 (for higher AP, still realtime on T4/Orin)

- Export via the repo's tools to ONNX, then `trtexec` as above (batch=1, 640 or 800).

### ROI classifier (Path B)

- Train a **multi-label** head (MobileNetV3 / EfficientNet-B0) on crops:
  - labels: ppe.hard\_hat , ppe.safety\_vest , tool (each  $\in \{0,1\}$ )
- Export ONNX: `torch.onnx.export` with `input: 1x3x224x224`
- Build TRT FP16: `trtexec --onnx=ppe_cls.onnx --saveEngine=ppe_cls-fp16.engine --fp16`

## 2.3 Detector service ( src/detectors/object/main.py )

### Responsibilities:

- Read frames (shared queue / RTSP reader)

- Preprocess (letterbox to 640×640, FP16)
- TRT enqueue → outputs
- Decode (logits → boxes, scores, classes, NMS)
- Publish NDJSON events

NDJSON event (Path A – trained PPE detector):

```
{"type":"object","camera_id":"cam1","ts":1730227201.123, "detections":
[{"bbox": [x,y,w,h], "cls":"hard_hat", "conf":0.93}]}
```

NDJSON (Path B – general detector):

```
{"type":"object","camera_id":"cam1","ts":1730227201.123, "detections":
[{"bbox": [x,y,w,h], "cls":"person", "conf":0.89}]}
```

### Minimal post-proc knobs

- score\_thresh : 0.25–0.4
- nms\_iou : 0.45
- max detections per frame: 100

## 2.4 Tracking ( src/trackers/bytetrack.py )

- Input: per-frame detections (x,y,w,h, score, class)
- Output: add track\_id with stable IDs
- Publish:

```
{"type":"object_tracked","camera_id":"cam1","ts":..., "detections":
[{"track_id":7,"bbox": [...], "cls":"person", "conf":0.89}]}
```

## 2.5 ROI classifier (Path B)

- Subscribe to object\_tracked
- For each person (or tool-like) box: crop ROI(s)
  - Head/torso crop: use bbox heuristics or a tiny head-finder
- Run TRT engine ppe\_cls-fp16.engine
- Attach attributes:

```
{"type":"object_attributed","camera_id":"cam1","ts":..., "detections":
[{"track_id":7,"bbox": [...], "cls":"person", "attributes":{"ppe":
```

```
{"hard_hat":true,"safety_vest":false},"tool":true}}}]}
```

## 3) Pose estimation

### 3.1 Choose your model

- RTMPose-m (best accuracy/speed balance, top-down)
- RTMO (one-stage, great for crowded scenes)
- YOLO-pose (easiest export path if you're already in Ultralytics)

### 3.2 Export → TensorRT

#### A) RTMPose via MMDeploy (canonical)

1. Convert to ONNX with the official MMPose/MMDeploy toolchain
2. `trtexec --onnx=rtmpose.onnx --saveEngine=rtmpose-fp16.engine --fp16 --shapes=input:1x3x640x640`

#### B) YOLOv8-pose (fastest path)

- `yolo pose export model=yolov8n-pose.pt format=onnx imgsz=640`
- `trtexec ... --saveEngine=yolov8n-pose-fp16.engine --fp16`

### 3.3 Pose service ( src/detectors/pose/main.py )

- Preprocess → TRT → decode heatmaps/offsets or direct kpts
- Output schema:

```
{"type":"pose","camera_id":"cam1","ts":..., "detections":[ {"bbox":
[x,y,w,h],"keypoints":[[x,y,conf],... 17], "conf":0.91, "person_id":12}
]}
```

Tracking for pose: reuse `object_tracked` IDs by associating each skeleton with the nearest tracked `person` `bbox`; fallback to simple KPT-IoU matching if needed.

### 3.4 Pose velocity analytics ( src/analytics/pose\_velocity.py )

- For each `person_id`, keep a ring buffer of last N frames (kpts)
- Velocity per joint  $v = ||p_t - p_{t-1}|| / \Delta t$
- Publish:

```
{"type":"pose_velocity","camera_id":"cam1","ts":..., "track_id":12,
"joints":{"wrist_r":0.42,"ankle_l":0.17}, "v_mean":0.21}
```

## 4) Analytics: PPE rules (works for Path A or B)

Config ( configs/analytics/ppe\_rules.yaml )

```
actor_class: person dwell_seconds: 3 required: hard_hat: true
safety_vest: true tool_allowed: false
```

Rule engine ( src/analytics/ppe\_rules.py )

- Subscribe to `object_tracked` (Path A: class already `hard_hat` , etc.) or `object_attributed` (Path B)
- Maintain per `track_id` timers for dwell
- Emit violations:

```
{"type":"alert","subtype":"missing_ppe","track_id":7,"camera_id":"cam1",
"missing":["hard_hat"],"dwell":3.2,"ts":...}
```

## 5) Event bus (file-based to start)

Write: append NDJSON to e.g. `artifacts/detections/object.ndjson`

Read: tail-like stream per consumer

Common pub helper:

```
def publish(path, obj): with open(path, "a") as f:
    f.write(json.dumps(obj) + "\n")
```

(When scaling, swap to Kafka/Redis without touching producers/consumers—just replace the adapter.)

## 6) Docker & Compose

## 6.1 Detector Dockerfiles (GPU)

src/detectors/object/docker/Dockerfile

```
FROM nvcr.io/nvidia/cuda:12.2.0-runtime-ubuntu20.04 RUN apt-get update
&& apt-get install -y python3-pip libgl1-mesa-glx libsm6 libxrender1
libxext6 && rm -rf /var/lib/apt/lists/* COPY requirements.txt
/app/requirements.txt RUN pip3 install --no-cache-dir -r
/app/requirements.txt WORKDIR /app COPY . /app ENV
TRT_LOGGER_SEVERITY=ERROR CMD ["python3",
"src/detectors/object/main.py", "--config",
"configs/detectors/object.yaml"]
```

(Use nvcr.io/nvidia/l4t-jetpack base on Jetson.)

## 6.2 docker-compose.yml

```
version: "3.8" services: detect_object: build: { context: .,
dockerfile: src/detectors/object/docker/Dockerfile } runtime: nvidia
environment: [ NVIDIA_VISIBLE_DEVICES=all ] volumes: -
./models:/app/models - ./configs:/app/configs -
./artifacts:/app/artifacts tracker: build: { context: ., dockerfile:
src/trackers/docker/Dockerfile } runtime: nvidia volumes: [
./artifacts:/app/artifacts ] detect_pose: build: { context: .,
dockerfile: src/detectors/pose/docker/Dockerfile } runtime: nvidia
volumes: [ ./models:/app/models, ./configs:/app/configs,
./artifacts:/app/artifacts ] analytics: build: { context: ., dockerfile:
src/analytics/docker/Dockerfile } volumes: [ ./configs:/app/configs,
./artifacts:/app/artifacts ]
```

Make sure host has `--gpus all` (Compose v2 uses `runtime: nvidia` or `deploy.resources.reservations.devices`).

## 7) Configs (detectors)

configs/detectors/object.yaml

```
engine: models/object/engines/yolo-fp16.engine imgsz: 640 score_thresh:
0.35 nms_iou: 0.45 max_det: 100 publish:
artifacts/detections/object.ndjson
```

configs/detectors/pose.yaml



```
engine: models/pose/engines/rtpose-fp16.engine imgsz: 640
score_thresh: 0.3 publish: artifacts/detections/pose.ndjson
```

configs/classifier/ppe.yaml (Path B)

```
engine: models/classifier/ppe_cls-fp16.engine input_size: 224 labels:
["ppe.hard_hat", "ppe.safety_vest", "tool"] thresholds: [0.5, 0.5, 0.6]
publish: artifacts/detections/object_attributed.ndjson
```

## 8) Tests (ship-blocking)

### 8.1 Unit

- **Postproc decode:** raw logits → 5 known boxes (golden)
- **NMS:** deterministic results for overlaps
- **Pose decode:** kpt heatmaps → 17 points within tolerance
- **Classifier:** logits → multi-label thresholding

### 8.2 Integration (short clips)

- `tests/e2e_object.yaml`: run `detect_object` (+ tracker) on a 10-sec clip  
→ assert at least one person with stable `track_id`  
→ if Path B, assert attributes JSON keys exist.
- `tests/e2e_pose.yaml`: run `detect_pose` on same clip  
→ assert 17 kpts per detected person  
→ if velocity enabled, assert `pose_velocity` events emitted.

## 9) Observability & ops

- **Health endpoints:** each service exposes `/:health` (HTTP 200)
- **Metrics:** Prometheus counters: FPS, latency (p50/p95), dropped frames, publishes/sec
- **Logs:** structured JSON; at least: `camera_id`, `frame_id`, `model`, `lat_ms`

## 10) Performance checklist

- Pre-allocate device buffers; reuse CUDA streams
- Convert to **NHWC/NCHW** as model expects once; avoid per-frame reallocations
- Use **FP16** first; INT8 later if you hit a wall (ensure representative calibration)
- Pin host memory for H2D copies; batch size = 1 for live streams
- Run **NMS on GPU** if available; otherwise optimize CPU NMS (NumPy → torch/onnx runtime kernels)
- Turn on **affinity**: dedicate cores to decode vs infer on Jetson

## 11) What to do today (P0 checklist)

1. **Pick path** for PPE:
  - If you **have labeled PPE data** → Path A (single-stage).
  - If you **don't yet** → Path B (two-stage with ROI classifier).
2. **Build engines**
  - Detector (YOLO or RT-DETRv2) → ONNX → TRT FP16
  - Pose (RTMPose or YOLO-pose) → ONNX → TRT FP16
  - (Path B) ROI classifier → ONNX → TRT
3. **Implement services**
  - `src/detectors/object/main.py` + `postproc.py`
  - `src/trackers/bytetrack.py`
  - (Path B) `src/classifier/ppe/infer.py`
  - `src/detectors/pose/main.py` + `pose_velocity.py`
4. **Wire configs & compose**
  - Fill `configs/*` as shown; mount `models/` & `artifacts/` volumes
5. **Run e2e tests** on short clips; fix thresholds until stable
6. **Add health + metrics**; set alerts on FPS < target or publish stalls

## 12) Minimal code stubs (drop-in)

## src/detectors/object/main.py (skeleton)

```

import json, time import numpy as np import tensorrt as trt,
pycuda.driver as cuda, pycuda.autotinit from postproc import decode def
load_engine(path): logger = trt.Logger(trt.Logger.ERROR) with open(path,
"rb") as f, trt.Runtime(logger) as rt: return
rt.deserialize_cuda_engine(f.read()) def allocate(engine): ctx =
engine.create_execution_context() bindings, host, device = [], {}, {}
for i, name in enumerate(engine): dtype =
trt.nptype(engine.get_binding_dtype(name)) shape =
tuple(engine.get_binding_shape(i)) size = np.prod(shape) host[name] =
cuda.pagelocked_empty(size, dtype) device[name] =
cuda.mem_alloc(host[name].nbytes) bindings.append(int(device[name]))
stream = cuda.Stream() return ctx, bindings, host, device, stream def
infer(ctx, bindings, host, device, stream, inp):
np.copyto(host["images"], inp.ravel())
cuda.memcpy_htod_async(device["images"], host["images"], stream)
ctx.execute_async_v2(bindings, stream.handle)
cuda.memcpy_dtoh_async(host["output0"], device["output0"], stream)
stream.synchronize() return host["output0"] def publish(path, obj): with
open(path, "a") as f: f.write(json.dumps(obj)+"\n") def main(cfg): eng =
load_engine(cfg["engine"]) ctx, bindings, host, device, stream =
allocate(eng) while True: frame = read_frame_somewhat() # BGR HxWx3 img,
scale, pad = preprocess_letterbox(frame, 640) out = infer(ctx, bindings,
host, device, stream, img) dets = decode(out, scale, pad,
cfg["score_thresh"], cfg["nms_iou"]) evt =
{"type": "object", "camera_id": cfg.get("camera_id", "cam1"),
"ts": time.time(), "detections": dets} publish(cfg["publish"], evt) if
__name__ == "__main__": import yaml, sys cfg =
yaml.safe_load(open(sys.argv[-1])) main(cfg)

```

*(Pose is analogous; postproc differs.)*

If you want, I can generate **ready-to-run**

**stubs** for detect\_object , detect\_pose , bytetrack , the ROI classifier, and sample configs exactly in your tree so you can docker compose up and see events flowing.